



Cygnacom
Solutions



Analyzing Block Device Timing Events as a Source of Entropy

by Kirill Sinitski and Mike Ounsworth



Random Numbers

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

Motivation

- Comment in Linux kernel's random.c says:

“add_disk_randomness() uses what amounts to the seek time .., as input to the entropy pool. **Note that high-speed solid state drives with very low seek times do not make for good sources of entropy**, as their seek times are usually fairly consistent.”

- Our work aims to provide data to validate / invalidate this statement.

About us

- FIPS and Common Criteria Labs
 - Accredited testing laboratories
 - NIAP, NIST, CSE
- Entrust Datacard
 - Identity and Access Management
 - ID and financial card printers
 - SSL certificates



Entrust Datacard™

Trusted Identities | Secure Transactions

Analyzing Block Device Timing Events as a Source of Entropy

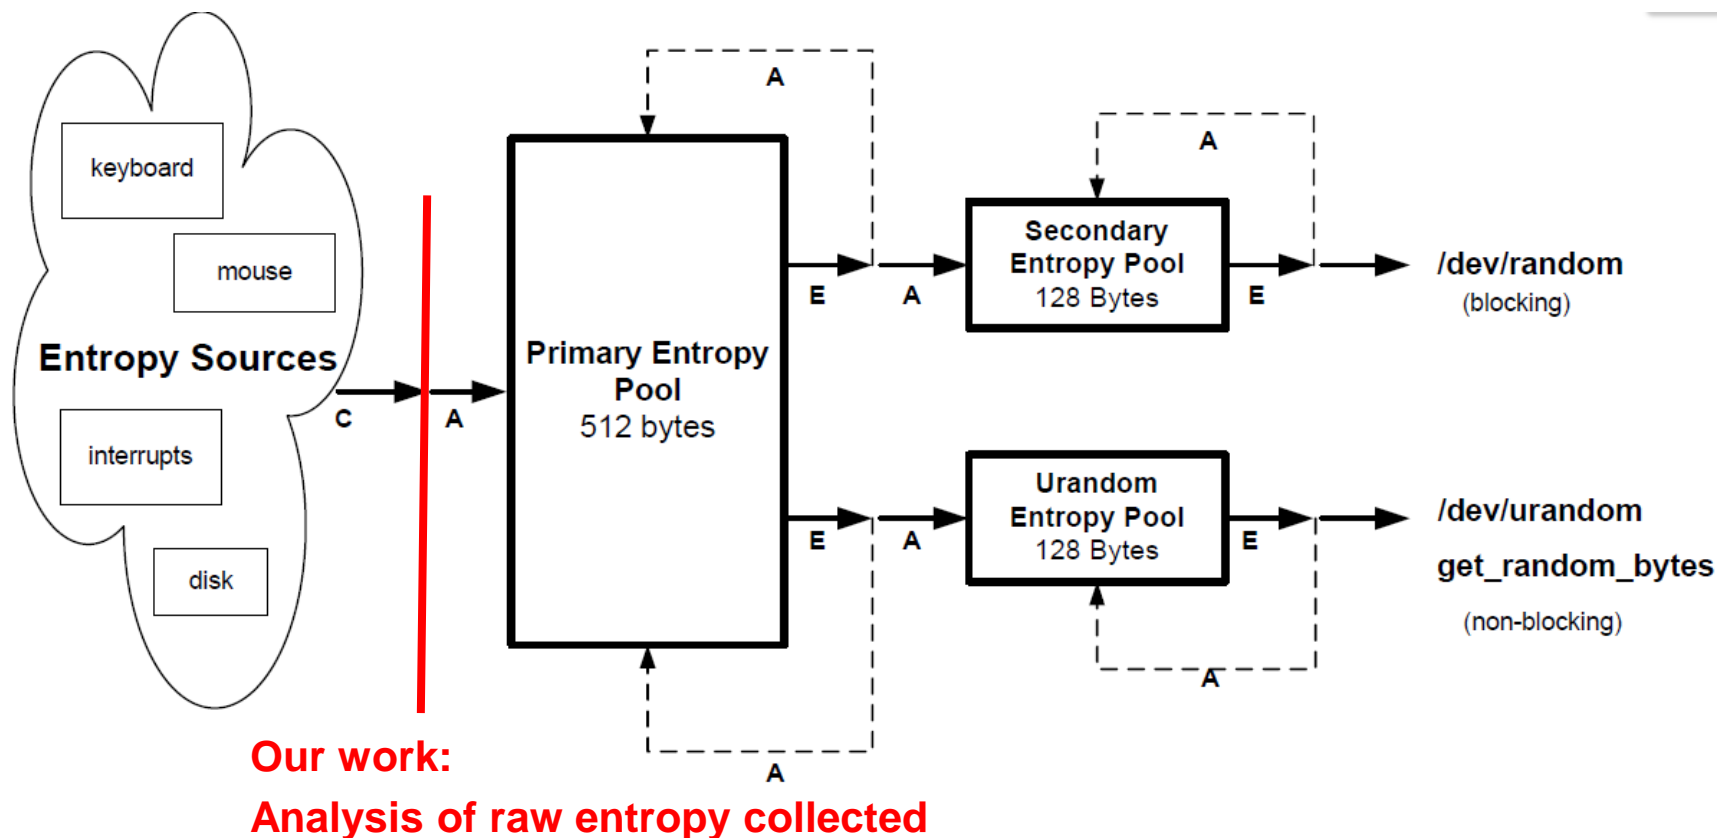
KERNEL

Entropy - Definitions

- **General:**
Lack of order or unpredictability.
- **Physics:**
A measure of disorder in the universe.
- **Information Theory:**
Uncertainty of a random variable.
- **Computing:**
The randomness collected by an operating system or application for use in cryptography or other uses that require unpredictable data.
 - Linux PRNG collects entropy from the timestamps of various events (user inputs, disk timings, and interrupt timings)

Linux Random Number Generator

<linux kernel>/drivers/char/random.c



Entropy sources for random.c (as of 4.16)

- `add_device_randomness(...)`
- `add_input_randomness(...)`
- `add_interrupt_randomness(...)`
- `add_disk_randomness(...)`
- `add_hwgenerator_randomness(...)`

Timing of events

- Jiffies
 - Kernel-level interrupt-based timer
 - Timer tick rate 'HZ' set at kernel compile time
 - Typically 100 – 250 Hz (4 ms to 10 ms intervals)
 - (mostly) Hardware agnostic
- CPU cycles
 - Hardware-specific
 - CPU clock cycles
 - Could be variable, depending on CPU features (e.g., dynamic frequency scaling)

What is inside an entropy event?

```
static void add_timer_randomness(struct timer_rand_state *state, unsigned num,
                                const char * source)
```


```
{
    struct {
        long jiffies;      ← timestamp of event in jiffies
        unsigned cycles;   ← timestamp of event in cycles
        unsigned num;      ← event type (keyboard scan code, interrupt code)
    } sample;
    ...
    mix_pool_bytes(&input_pool, &sample, sizeof(sample), NULL);
    ...
    credit_entropy_bits(&input_pool, min_t(int, fls(delta>>1), 11));
}
```

Sample event: `jiffies: 4294671006, cycles: 2129478438, num: ide1`
 `entropy_credited: 2`

credit_entropy_bits(..)

Linux Entropy Estimator

- How many bits of randomness in a given timestamp sample t_n ?
- Polynomial interpolation – 1st, 2nd, 3rd Differences



MIN_n - the minimum D_i value for the newly added sample t_n .

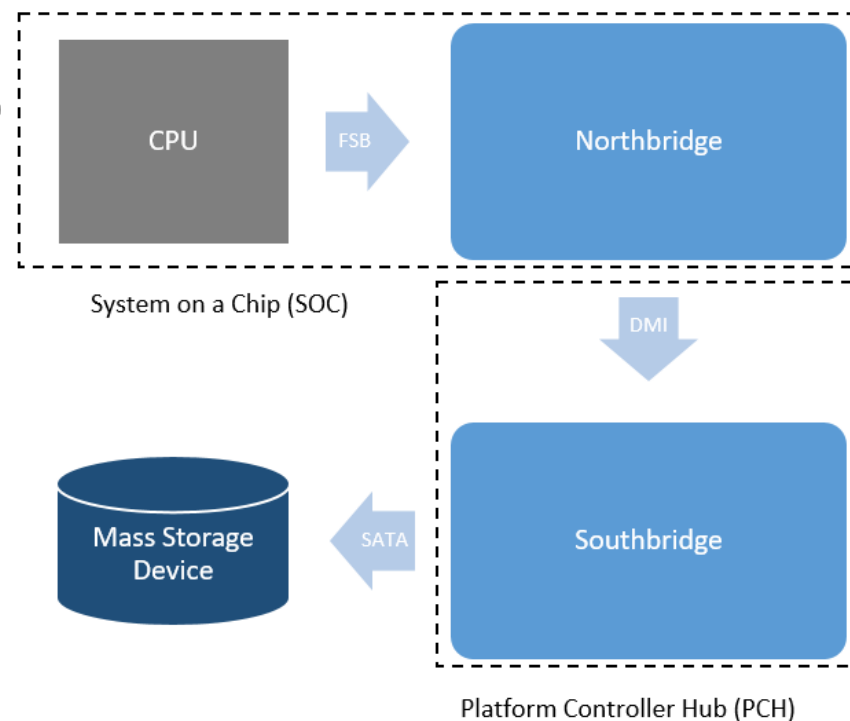
$$ent_n = \begin{cases} 0 & \text{if } MIN_n \leq 1 \\ \left\lfloor \log_2 (MIN_n) \right\rfloor & \text{if } 2 \leq MIN_n \leq 2^{12} \\ 11 & \text{otherwise} \end{cases}$$

Analyzing Block Device Timing Events as a Source of Entropy

FIRMWARE & CONTROLLERS

Hardware Architecture

- Hub Architecture (HA)
- Platform Controller Hub (PCH)
- System on a Chip (SoC)



- Test platform – Intel C610/X99 ‘Wellsburg’ PCH
 - 20 Gbit/s DMI 2.0, SATA 3.0 controller

Interfaces

- Storage Device Interfaces
 - AHCI vs Parallel ATA (IDE emulation)
 - AHCI HBA (host bus adapter) DMA (direct memory access)
- SATA
 - Native Command Queuing (NCQ)
 - Tagged Command Queuing (TCQ)
 - SATA 3.2 at 1969 MB/s, ATA-3 at 33 MB/s
- NVMe
 - Non-volatile Memory Express (NVMe)
 - PCIe and M.2

Boot Process

- Basic Input/Output System (BIOS)
 - BIOS executes Master Boot Record (MBR)
 - MBR executes Grand Unified Bootloader (GRUB)
 - GRUB executes start_kernel()
- Universal Extensible Firmware Interface (UEFI)
 - Boot Services load EFI System Partition (ESP)
 - Kernel boot stub (acting as an EFI application)
 - EFI handover to GRUB2
 - GRUB2 executes start_kernel()

Analyzing Block Device Timing Events as a Source of Entropy

HARDWARE

Test Platform

- Test platform – Dell Power Edge T430
 - Intel C610/X99 with Intel Xeon E5-2603v3
 - 20 Gbit/s DMI 2.0, integrated SATA 3.0 controller

Hard Drives and Performance

- Seagate Hybrid Drive ST2000DX001 2TB
 - 210 MB/s, 4.2ms seek
 - 7200 RPM, 64MB Cache
- WD VelociRaptor WD1500HLFS 150GB
 - 126 MB/s, 4.6ms seek
 - 10000 RPM, 16MB Cache
- Seagate Cheetah 15K.5 1TB
 - 128 MB/s, 6.6ms seek
 - 15000 RPM, 16MB Cache
- Western Digital Black 2TB
 - 164 MB/s, 12.0ms seek
 - 7200 RPM, 64 MB Cache

Solid State Drives and Performance

- Intel 540S Series 480GB
 - 560 MB/s, 0.05 ms access
- Intel 335 Series 180GB
 - 500 MB/s, 0.2 ms access
- WD SiliconEdge Blue 128GB
 - 250 MB/s, 0.25 ms access

Analyzing Block Device Timing Events as a Source of Entropy

TEST HARNESS

Log each entropy event to kernel logs

```
random.c :: add_timer_randomness( .. )
{
    ...
    mix_pool_bytes(&input_pool, &sample, sizeof(sample), NULL);
    ...
    credit_entropy_bits(&input_pool, min_t(int, fls(delta>>1), 11));

    // Test Harness - log raw data
    printk(KERN_DEBUG, source, input_pool.entropy_total,
           input_pool.entropy_count, sample.jiffies,
           sample.cycles, min_t(int, fls(delta>>1), 11) );
}
```

Typical journalctl -k log entry:

```
Jan 12 09:56:24 localhost kernel: DISK_RANDOMNESS: ent_total: 42,
ent_count: 244, jiffies: 4294671331, cycles: 2364333088, min_t: 2
```

Experimental setup

Rationale: wanted to best represent entropy collection in a datacenter environment.

Setup:

- Dell PowerEdge T430 with Intel Xeon E5-2603v3
- CentOS 7.2 Minimal
- Kernel 3.10.0-327 (from centos.org mirror)
 - Re-compiled using default options to insert our test harness into random.c
- Continuous reboot cycles
 - Until 1.5 million events or 3 days elapsed (~ 400 – 3,000 reboots)

reboot.sh

```
...
# log data
journalctl -k -p 7 | grep RANDOMNESS >> $DATA_FILE
...
if (( ELAPSED_TIME > EXPERIMENT_MAX_RUNTIME
      -o `wc -l < $DATA_FILE` > NUM_SAMPLES_TO_COLLECT ))
then
    echo "Ending experiment"
else
    reboot
fi
```

Analyzing Block Device Timing Events as a Source of Entropy

DATA

Data

- Analyzed
 - Cycles and Jiffies extracted and processed
 - Statistical analysis of 6 bit and 8 bit sample with NIST tool
- Collected, not analyzed
 - Internal entropy estimate per sample is extracted
 - Internal Linux estimate vs. post-hoc statistical analysis
 - Each reboot is logged with a timestamp
 - Variability between individual reboots
 - State of entropy pool at various stages during the boot cycle
 - Other possible research ...

Set it free

- All our data, harness, scripts, and procedures are available on Github:

<https://github.com/ounsworth/LinuxEntropyAnalysis>


Analyzing Block Device Timing Events as a Source of Entropy

RESULTS

Post-hoc Entropy Analysis

- NIST python tool
 - Collision
 - Partial collection
 - Markov
 - Compression
 - Frequency
- Min Entropy
 - Lowest of the test scores


BIOS vs. UEFI




Cycles vs. Jiffies

- Min Entropy

Cycles vs Jiffies Min Entropy



Cycles vs Jiffies Markov

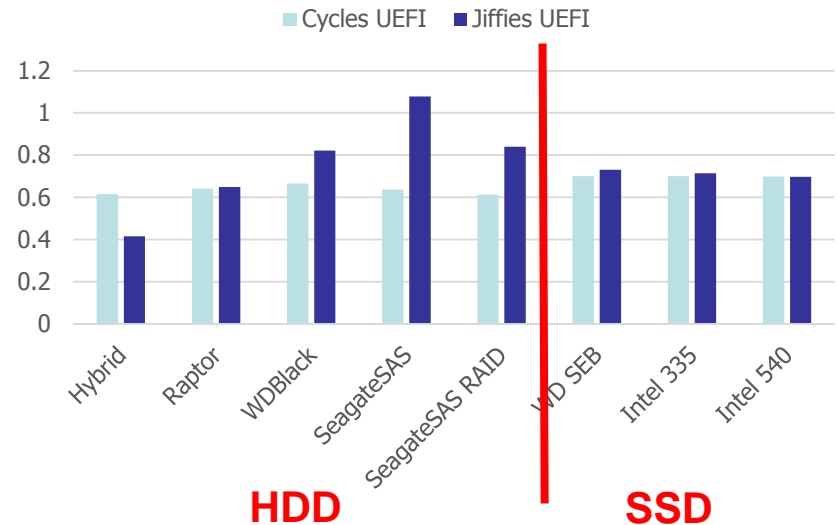


- Markov


HDD vs SSD

- Min Entropy

Cycles vs Jiffies Min Entropy



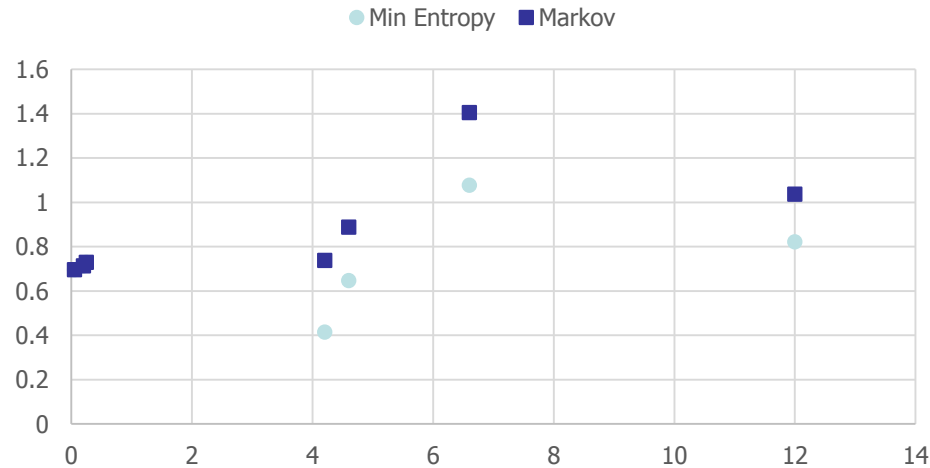
Cycles vs Jiffies Markov



- Markov

Hardware Considerations

Access Time vs. Entropy



- Published Access Time/Latency numbers were not verified

Summary of Findings


- BIOS vs UEFI are comparable for entropy generation
- Cycles vs Jiffies are comparable per non-duplicate event entropy using min-entropy
- Both SSD and HDD generate comparable entropy per event
- HDD generates 15x as many entropy events per reboot cycle as SSD
- HDDs generate ~50% duplicate jiffie events
- SSDs generate ~1% duplicate jiffie events

Speculation

- Speculation (noun) - the forming of a theory or conjecture without firm evidence
- SATA SSDs generate good per event entropy, just less events than HDDs
- Markov is more discerning test than min-entropy
- Disk drive speed and access times had no detectable effect on per event entropy within the same system, controller performance probably determines entropy
- NVMe parallelism has a potential to drastically change SDD performance as a source of entropy

Entropy Assessment Report

- Part of FIPS and CC certifications
 - Recommendation for Random Number Generation Using Deterministic Random Bit Generators, NIST SP 800-90A
 - Recommendation for the Entropy Sources Used for Random Bit Generation, NIST SP 800-90B
- Entropy Assessment Report (EAR)
 - Per source data analysis
 - Operating Conditions
 - Health Testing
 - Description of conditioning function
 - Description of DRBG seed creation



Questions?

